Create and Manage Cloud SQL for PostgreSQL Instances: Challenge Lab

Task 1. Migrate a stand-alone PostgreSQL database to a Cloud SQL for PostgreSQL instance

In this task you must migrate the stand-alone PostgreSQL orders database running on the postgres-vm virtual machine to a Cloud SQL for PostgreSQL instance using a Database Migration Services continuous migration job and VPC Peering connectivity.

Prepare the stand-alone PostgreSQL database for migration

In this sub-task you must prepare the stand-alone PostgreSQL database so that it satisfies the requirements for migration by Database Migration Services.

To complete this sub-task you must complete the following steps:

1. Enable the Google Cloud APIs required for Database Migration Services.

Database Migration Services require the **Database Migration API** and the **Service Networking API** to be enabled in order to function. You must enable these APIs for your project.

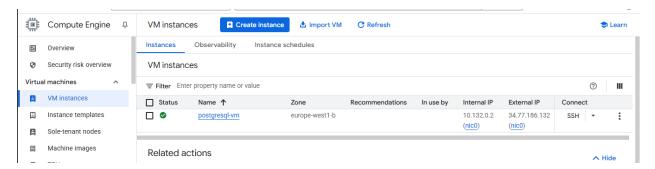
- 2. Upgrade the target databases on the postgres-vm virtual machine with the pglogical database extension.
- 3. You must install and configure the **pglogical** database extension on the stand-alone PostgreSQL database on the postgres-vm Compute Instance VM. The pglogical database extension package that you must install is named postgresql-13-pglogical.
- 4. To complete the configuration of the **pglogical** database extension you must edit the PostgreSQL configuration file /etc/postgresql/13/main/postgresql.conf to enable the **pglogical** database extension and you must edit the /etc/postgresql/13/main/pg_hba.conf to allow access from all hosts.
- 5. Create a dedicated user for database migration on the stand-alone database.
- 6. The new user that you create on the stand-alone PostgreSQL installation on the postgres-vm virtual machine must be configured using the following user name and password:
- Migration user name: Postgres Migration User

- Migration user password : DMS_1s_cool!
- 7. Grant that user the required privileges and permissions for databases to be migrated.

Database Migration Services require that the migration user has privileges to specific schemata and relations of the target databases for migration, in this case that is the orders database.

The Database Migration Service requires all tables to be migrated to have a primary key.

- 8. You must make sure that all of the tables in the orders database have a primary key set before you start the migration.
- distribution_centers
- inventory_items
- order_items
- products
- users



Migrate the stand-alone PostgreSQL database to a Cloud SQL for PostgreSQL instance In this sub-task you must perform the migration using Database Migration Services.

To complete this sub-task you must complete the following steps:

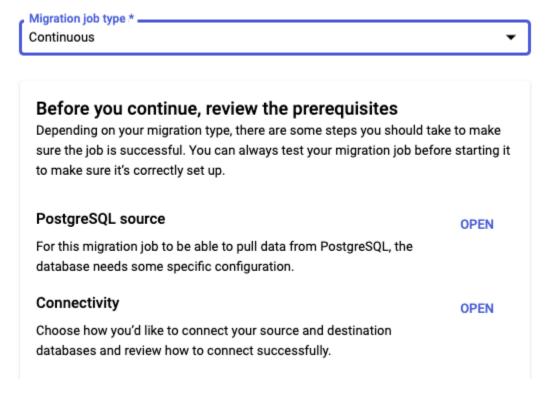
 Create a new Database Migration Service connection profile for the stand-alone PostgreSQL database, using the credentials of the Postgres Migration User migration user you created earlier.

• **Username**: Postgres Migration User

Password : DMS_1s_cool!

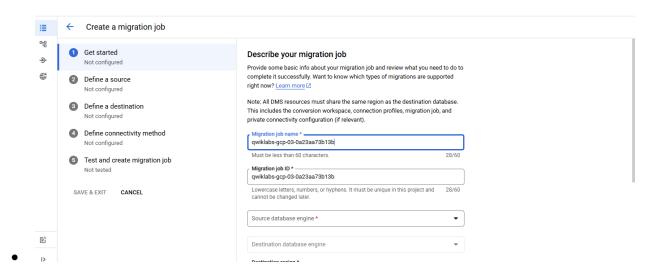
- Region must be set to REGION
- 2. You must configure the connection profile using the internal ip-address of the source compute instance.
- 3. Create a new continuous Database Migration Service job.

As part of the migration job configuration, make sure that you specify the following properties for the destination Cloud SQL instance:



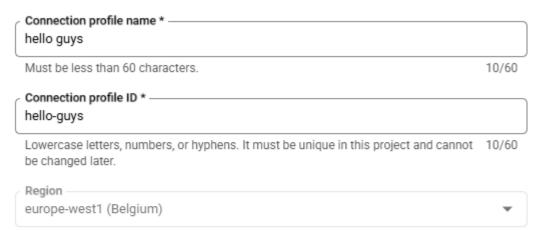
- The **Destination Instance ID** must be set to Migrated Cloud SQL for PostgreSQL Instance ID
- The **Password** for the migrated instance must be set to supersecret!
- Database version must be set to Cloud SQL for PostgreSQL 13
- Region must be set to REGION
- For Connections both Public IP and Private IP must be set.
- For Private IP use an automatically allocated IP range.
- Select **Enterprise** Cloud SQL edition.
- Select a dedicated core machine shapes with 2 vCPU and 8GB of memory.

- For the **Connectivity Method**, you must use **VPC peering** with the **default** VPC network.
- Test and then start the continuous migration job.



Create connection profile

Provide high-level information, specify a region, and enter connection details for your connection profile.



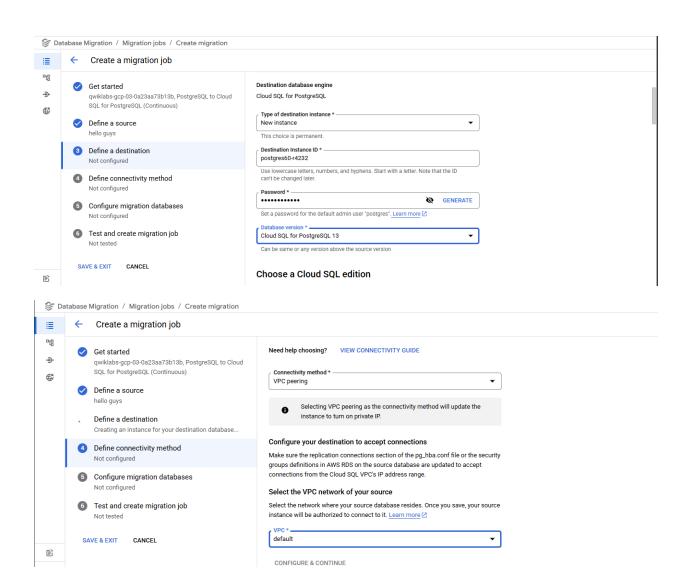
Define connection configurations

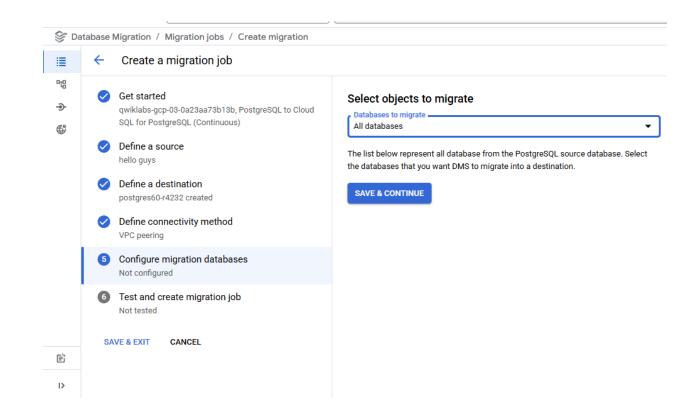
The combination of source and destination engines determines further configuration of your source connection profile.

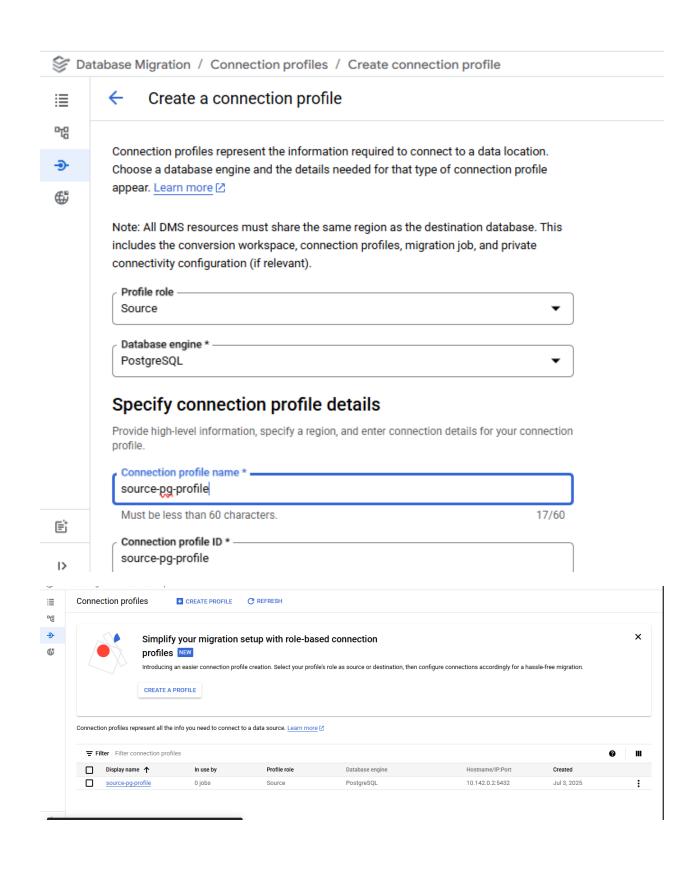


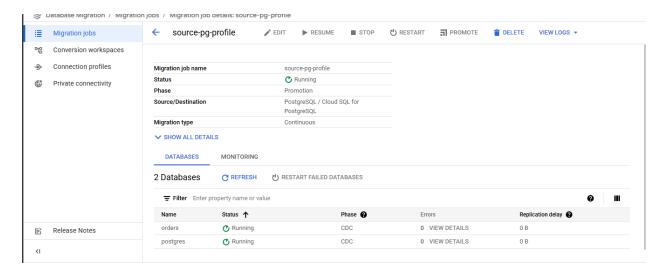


CANCE



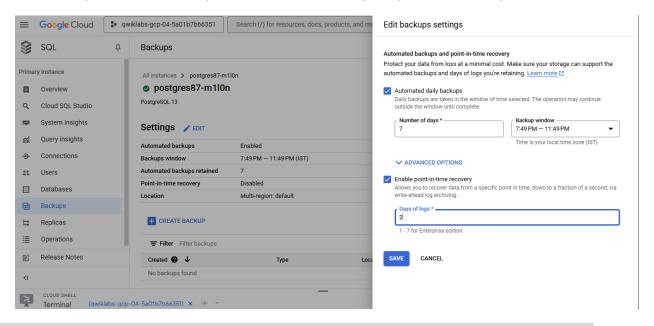






Task 2. Promote a Cloud SQL to be a stand-alone instance for reading and writing data

- In this task, you must complete the migration by promoting the Cloud SQL for PostgreSQL instance to a stand-alone instance.
- When the promotion is complete, the status of the job updates to Completed.



Task 3. Implement Cloud SQL for PostgreSQL IAM database authentication

In this task you must configure the newly migrated Cloud SQL for PostgreSQL instance to support Cloud IAM users and IAM database authentication. You are also required to patch the Cloud SQL for PostgreSQL instance, to add the public ip-address of the postgres-vm virtual machine to the list of networks that are allowed to connect to the instance.

To complete this task you must complete the following steps:

- 1. Patch the Migrated Cloud SQL for PostgreSQL Instance ID Cloud SQL instance to allow connections from the public ip-address of the postgres-vm virtual machine.
 - In the Migrated Cloud SQL for PostgreSQL Instance ID Cloud SQL instance,
 go to connections > Networking.
 - b. Under the Public IP, click on **ADD A NETWORK**. For the network, use the external IP of the postgres-vm virtual machine.
- In the Migrated Cloud SQL for PostgreSQL Instance ID Cloud SQL instance, create a Cloud SQL IAM user using the lab student ID, Qwiklabs user account name, as the principal account name.
 - a. Click Users > Add user account, then select Cloud IAM.
- 3. Grant SELECT permission to the Cloud IAM user for the orders table.
 - a. In the Migrated Cloud SQL for PostgreSQL Instance ID Cloud SQL instance, go to Overview. Under Connect to this instance, click on Open Cloud Shell.
 - b. For the password enter supersecret!. Then connect to the orders database using \c orders; command.
 - c. Again for the password enter supersecret!.
 - d. Use the following command to grant SELECT permission. Replace the **Table_Name** and **Qwiklabs_User_Account_Name** variables with the correct values.

GRANT SELECT ON Table_Name TO "Qwiklabs_User_Account_Name";

- 1. Run the following query as the Qwiklabs user account name user in the migrated database to confirm that the Qwiklabs user account name can select data from the orders table.
- 2. SELECT COUNT(*) FROM orders

Task 4. Configure and test point-in-time recovery

In this task you must configure point-in-time recovery on a Cloud SQL for PostgreSQL instance and then test it by using point-in-time recovery to create a cloned instance at a point in time that rolls back some changes.

To complete this task you must complete the following steps:

- 1. Enable backups on the Cloud SQL for PostgreSQL instance.
 - a. In the Migrated Cloud SQL for PostgreSQL Instance ID Cloud SQL instance, go to **Overview**. Click on **edit** > **Data Protection**.
 - b. Enable point-in-time recovery and set the number of retained transaction log days to Point-in-time recovery retention days.
- 2. Make a note of the timestamp for the point-in-time you wish to restore to using the following command.
- 3. Make some changes to the database after this timestamp.
 - a. In the Migrated Cloud SQL for PostgreSQL Instance ID Cloud SQL instance, go to Overview. Under Connect to this instance, click on Open Cloud Shell.
 - b. For the password enter supersecret!. Then connect to the orders database using \c orders; command.
 - c. Again for the password enter supersecret!.
 - d. You must add a row of data to the orders.distribution_centers table.
- 4. Use point-in-time recovery to create a clone that replicates the instance state at your chosen timestamp.
 - a. Use the following command to clone the instance by replacing the CLOUDSQL_INSTANCE, NEW_INSTANCE_NAME and TIME_STAMP variables with the correct values.

gcloud sql instances clone \$CLOUDSQL_INSTANCE \$NEW_INSTANCE_NAME \
--point-in-time \$TIME_STAMP

b. For the **Cloned instance name** you must use the name postgres-orders-pitr.

The new instance will not be used, but do not discard it, as it will be required to confirm that you have correctly completed the lab.

